

E1 - Análise de dados e Regressão Linear: Preço de imóveis

Considere a base de dados "Ames Housing Dataset". Com o objetivo de desenvolver um modelo de predição do preço do imóvel, desenvolva os itens a seguir e entregue a análise em arquivo do tipo powerpoint ou pdf

- Importando as bibliotecas

```
import pandas as pd
import numpy as np
import statsmodels
import seaborn
from matplotlib import pyplot as plt
pd.options.display.max_columns = 100
```

```
/Users/karinseeder/anaconda3/lib/python3.7/site-packages/pandas/compat/_optional.py:138: UserWarning: Pandas requires version '2.7.0' or newer of 'numexpr' (version '2.6.8' currently installed).
  warnings.warn(msg, UserWarning)
```

- Importando a base de dados

```
df = pd.read_csv('base_1ah.csv')
print(df.shape)
```

```
(1460, 81)
```

- Visualizando uma amostra da base_1ah

```
df.head()
```

	Id	MSSubClass	MSZoning	LotFrontage	LotArea	Street	Alley	LotShape	LandContour	Utilities	LotC
0	1	60	RL	65.0	8450	Pave	NaN	Reg	Lvl	AllPub	Insid
1	2	20	RL	80.0	9600	Pave	NaN	Reg	Lvl	AllPub	FR2
2	3	60	RL	68.0	11250	Pave	NaN	IR1	Lvl	AllPub	Insid
3	4	70	RL	60.0	9550	Pave	NaN	IR1	Lvl	AllPub	Corn
4	5	60	RL	84.0	14260	Pave	NaN	IR1	Lvl	AllPub	FR2

```
df = df.set_index('Id')
```

1 - Análise descritiva de variáveis

1.1 Estatísticas descritivas: frequência, proporção, média (\bar{x}), desvio padrão (s), quartis ($Q1$, \bar{x} , $Q3$) (1,0)

Avaliando o tipo das variáveis na base_1ah

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 1460 entries, 1 to 1460
Data columns (total 80 columns):
#   Column          Non-Null Count  Dtype
---  -
0   MSSubClass      1460 non-null  int64
1   MSZoning        1460 non-null  object
2   LotFrontage     1201 non-null  float64
3   LotArea        1460 non-null  int64
4   Street         1460 non-null  object
5   Alley          91 non-null    object
6   LotShape       1460 non-null  object
```

```
7 LandContour 1460 non-null object
8 Utilities 1460 non-null object
9 LotConfig 1460 non-null object
10 LandSlope 1460 non-null object
11 Neighborhood 1460 non-null object
12 Condition1 1460 non-null object
13 Condition2 1460 non-null object
14 BldgType 1460 non-null object
15 HouseStyle 1460 non-null object
16 OverallQual 1460 non-null int64
17 OverallCond 1460 non-null int64
18 YearBuilt 1460 non-null int64
19 YearRemodAdd 1460 non-null int64
20 RoofStyle 1460 non-null object
21 RoofMatl 1460 non-null object
22 Exterior1st 1460 non-null object
23 Exterior2nd 1460 non-null object
24 MasVnrType 1452 non-null object
25 MasVnrArea 1452 non-null float64
26 ExterQual 1460 non-null object
27 ExterCond 1460 non-null object
28 Foundation 1460 non-null object
29 BsmtQual 1423 non-null object
30 BsmtCond 1423 non-null object
31 BsmtExposure 1422 non-null object
32 BsmtFinType1 1423 non-null object
33 BsmtFinSF1 1460 non-null int64
34 BsmtFinType2 1422 non-null object
35 BsmtFinSF2 1460 non-null int64
36 BsmtUnfSF 1460 non-null int64
37 TotalBsmtSF 1460 non-null int64
38 Heating 1460 non-null object
39 HeatingQC 1460 non-null object
40 CentralAir 1460 non-null object
41 Electrical 1459 non-null object
42 1stFlrSF 1460 non-null int64
43 2ndFlrSF 1460 non-null int64
44 LowQualFinSF 1460 non-null int64
45 GrLivArea 1460 non-null int64
46 BsmtFullBath 1460 non-null int64
47 BsmtHalfBath 1460 non-null int64
48 FullBath 1460 non-null int64
49 HalfBath 1460 non-null int64
50 BedroomAbvGr 1460 non-null int64
51 KitchenAbvGr 1460 non-null int64
52 KitchenQual 1460 non-null object
53 TotRmsAbvGrd 1460 non-null int64
54 Functional 1460 non-null object
55 Fireplaces 1460 non-null int64
56 FireplaceQu 770 non-null object
57 GarageType 1379 non-null object
58 GarageYrBlt 1379 non-null float64
59 GarageFinish 1379 non-null object
60 GarageCars 1460 non-null int64
61 GarageArea 1460 non-null int64
62 GarageQual 1379 non-null object
63 GarageCond 1379 non-null object
64 PavedDrive 1460 non-null object
65 WoodDeckSF 1460 non-null int64
66 OpenPorchSF 1460 non-null int64
67 EnclosedPorch 1460 non-null int64
68 3SsnPorch 1460 non-null int64
69 ScreenPorch 1460 non-null int64
70 PoolArea 1460 non-null int64
71 PoolQC 7 non-null object
72 Fence 281 non-null object
73 MiscFeature 54 non-null object
74 MiscVal 1460 non-null int64
75 MoSold 1460 non-null int64
76 YrSold 1460 non-null int64
77 SaleType 1460 non-null object
78 SaleCondition 1460 non-null object
79 SalePrice 1460 non-null int64
dtypes: float64(3), int64(34), object(43)
memory usage: 923.9+ KB
```

Análise descritiva para variáveis numéricas

```
df.describe()
```

	MSSubClass	LotFrontage	LotArea	OverallQual	OverallCond	YearBuilt	YearRemodAdd	Ma
count	1460.000000	1201.000000	1460.000000	1460.000000	1460.000000	1460.000000	1460.000000	1452

	MSSubClass	LotFrontage	LotArea	OverallQual	OverallCond	YearBuilt	YearRemodAdd	Ma
mean	56.897260	70.049958	10516.828082	6.099315	5.575342	1971.267808	1984.865753	103.6
std	42.300571	24.284752	9981.264932	1.382997	1.112799	30.202904	20.645407	181.0
min	20.000000	21.000000	1300.000000	1.000000	1.000000	1872.000000	1950.000000	0.00
25%	20.000000	59.000000	7553.500000	5.000000	5.000000	1954.000000	1967.000000	0.00
50%	50.000000	69.000000	9478.500000	6.000000	5.000000	1973.000000	1994.000000	0.00
75%	70.000000	80.000000	11601.500000	7.000000	6.000000	2000.000000	2004.000000	166.0
max	190.000000	313.000000	215245.000000	10.000000	9.000000	2010.000000	2010.000000	1600

Análise descritiva das variáveis categóricas

```
df.describe(include=object)
```

	MSZoning	Street	Alley	LotShape	LandContour	Utilities	LotConfig	LandSlope	Neighborhood	Co
count	1460	1460	91	1460	1460	1460	1460	1460	1460	146
unique	5	2	2	4	4	2	5	3	25	9
top	RL	Pave	Grvl	Reg	Lvl	AllPub	Inside	Gtl	NAmes	Noi
freq	1151	1454	50	925	1311	1459	1052	1382	225	126

```
for i, x in enumerate(df.dtypes):
    if x == 'object':
        print(pd.crosstab(index=df[df.columns[i]], columns='freq', dropna=False))
        print('')
```

col_0	freq
MSZoning	
C (all)	10
FV	65
RH	16
RL	1151
RM	218
col_0	freq
Street	
Grvl	6
Pave	1454
col_0	freq
Alley	
Grvl	50
Pave	41
col_0	freq
LotShape	
IR1	484
IR2	41
IR3	10
Reg	925
col_0	freq
LandContour	
Bnk	63
HLS	50
Low	36
Lvl	1311
col_0	freq
Utilities	
AllPub	1459
NoSeWa	1

col_0	freq
LotConfig	
Corner	263
CulDSac	94
FR2	47
FR3	4
Inside	1052

col_0	freq
LandSlope	
Gtl	1382
Mod	65
Sev	13

col_0	freq
Neighborhood	
Blmngtn	17
Blueste	2
BrDale	16
BrkSide	58
ClearCr	28
CollgCr	150
Crawfor	51
Edwards	100
Gilbert	79
IDOTRR	37
MeadowV	17
Mitchel	49
NAmes	225
NPkVill	9
NWAmes	73
NoRidge	41
NridgHt	77
OldTown	113
SWISU	25
Sawyer	74
SawyerW	59
Somerst	86
StoneBr	25
Timber	38
Veenker	11

col_0	freq
Condition1	
Artery	48
Feedr	81
Norm	1260
PosA	8
PosN	19
RR Ae	11
RR An	26
RR Ne	2
RR Nn	5

col_0	freq
Condition2	
Artery	2
Feedr	6
Norm	1445
PosA	1
PosN	2
RR Ae	1
RR An	1
RR Nn	2

col_0	freq
BldgType	
1Fam	1220
2fmCon	31
Duplex	52
Twnhs	43
TwnhsE	114

col_0	freq
HouseStyle	
1.5Fin	154
1.5Unf	14
1Story	726
2.5Fin	8
2.5Unf	11
2Story	445
SFoyer	37
SLvl	65

col_0	freq
RoofStyle	
Flat	13

Gable	1141
Gambrel	11
Hip	286
Mansard	7
Shed	2

col_0	freq
RoofMatl	
ClyTile	1
CompShg	1434
Membran	1
Metal	1
Roll	1
Tar&Grv	11
WdShake	5
WdShngl	6

col_0	freq
Exterior1st	
AsbShng	20
AsphShn	1
BrkComm	2
BrkFace	50
CBlock	1
CemntBd	61
HdBoard	222
ImStucc	1
MetalSd	220
Plywood	108
Stone	2
Stucco	25
VinylSd	515
Wd Sdng	206
WdShing	26

col_0	freq
Exterior2nd	
AsbShng	20
AsphShn	3
Brk Cmn	7
BrkFace	25
CBlock	1
CmentBd	60
HdBoard	207
ImStucc	10
MetalSd	214
Other	1
Plywood	142
Stone	5
Stucco	26
VinylSd	504
Wd Sdng	197
Wd Shng	38

col_0	freq
MasVnrType	
BrkCmn	15
BrkFace	445
None	864
Stone	128

col_0	freq
ExterQual	
Ex	52
Fa	14
Gd	488
TA	906

col_0	freq
ExterCond	
Ex	3
Fa	28
Gd	146
Po	1
TA	1282

col_0	freq
Foundation	
BrkTil	146
CBlock	634
PConc	647
Slab	24
Stone	6
Wood	3

col_0	freq
BsmtQual	

Ex	121
Fa	35
Gd	618
TA	649

col_0	freq
BsmtCond	
Fa	45
Gd	65
Po	2
TA	1311

col_0	freq
BsmtExposure	
Av	221
Gd	134
Mn	114
No	953

col_0	freq
BsmtFinType1	
ALQ	220
BLQ	148
GLQ	418
LwQ	74
Rec	133
Unf	430

col_0	freq
BsmtFinType2	
ALQ	19
BLQ	33
GLQ	14
LwQ	46
Rec	54
Unf	1256

col_0	freq
Heating	
Floor	1
GasA	1428
GasW	18
Grav	7
OthW	2
Wall	4

col_0	freq
HeatingQC	
Ex	741
Fa	49
Gd	241
Po	1
TA	428

col_0	freq
CentralAir	
N	95
Y	1365

col_0	freq
Electrical	
FuseA	94
FuseF	27
FuseP	3
Mix	1
SBrkr	1334

col_0	freq
KitchenQual	
Ex	100
Fa	39
Gd	586
TA	735

col_0	freq
Functional	
Maj1	14
Maj2	5
Min1	31
Min2	34
Mod	15
Sev	1
Typ	1360

col_0	freq
FireplaceQu	
Ex	24

```

Fa          33
Gd          380
Po          20
TA          313

col_0      freq
GarageType
2Types      6
Attchd      870
Basement    19
BuiltIn     88
CarPort     9
Detchd      387

col_0      freq
GarageFinish
Fin         352
RFn         422
Unf         605

col_0      freq
GarageQual
Ex          3
Fa          48
Gd          14
Po          3
TA          1311

col_0      freq
GarageCond
Ex          2
Fa          35
Gd          9
Po          7
TA          1326

col_0      freq
PavedDrive
N           90
P           30
Y           1340

col_0      freq
PoolQC
Ex          2
Fa          2
Gd          3

col_0      freq
Fence
GdPrv      59
GdWo       54
MnPrv      157
MnWw       11

col_0      freq
MiscFeature
Gar2        2
Othr        2
Shed        49
TenC        1

col_0      freq
SaleType
COD         43
CWD         4
Con         2
ConLD       9
ConLI       5
ConLw       5
New        122
Oth         3
WD         1267

col_0      freq
SaleCondition
Abnorml    101
AdjLand     4
Alloca     12
Family     20
Normal     1198
Partial    125

```

1.2 Gráficos como: Gráficos de colunas, BoxPlot, dispersão (2,0)

Para gráficos numéricos plotamos boxplot e para categóricos gráficos de barra.




```


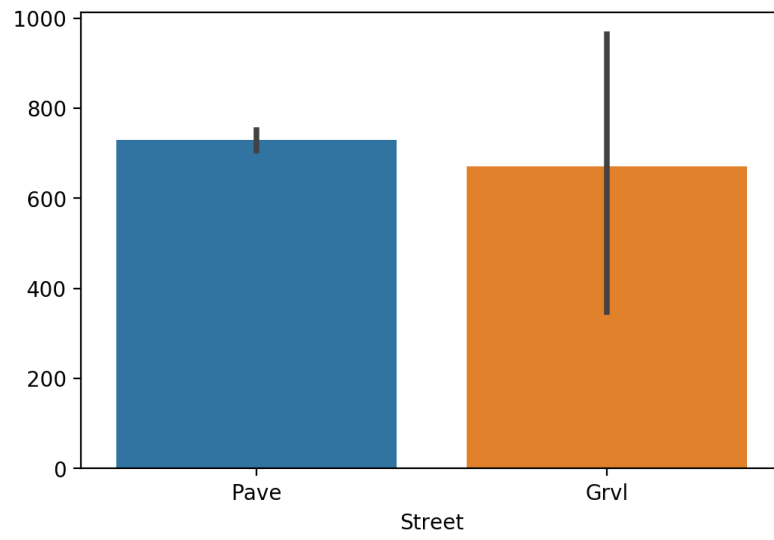

for i, x in enumerate(df.dtypes):
    if x == 'int64' or x == 'float64':
        plt.figure(i)
        seaborn.boxplot(y = df[df.columns[i]])
    elif x == 'int64' or x == 'object':
        plt.figure(i)
        seaborn.barplot(x = df[df.columns[i]], y = range(0, len(df)))

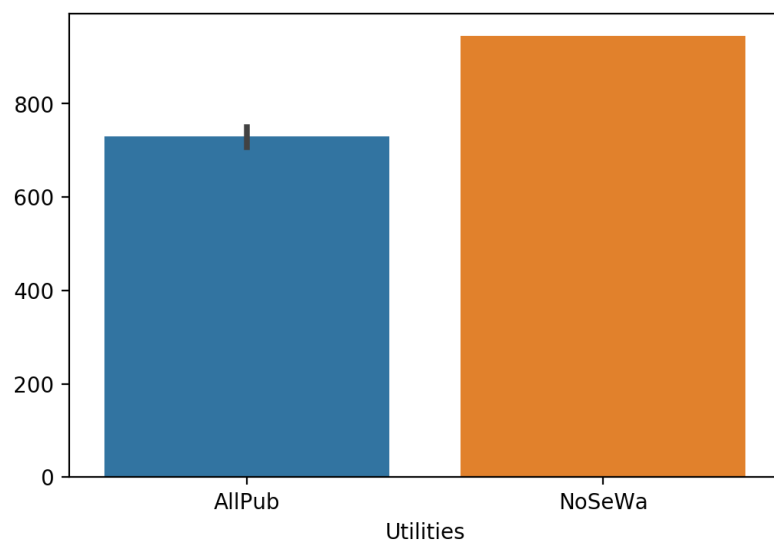
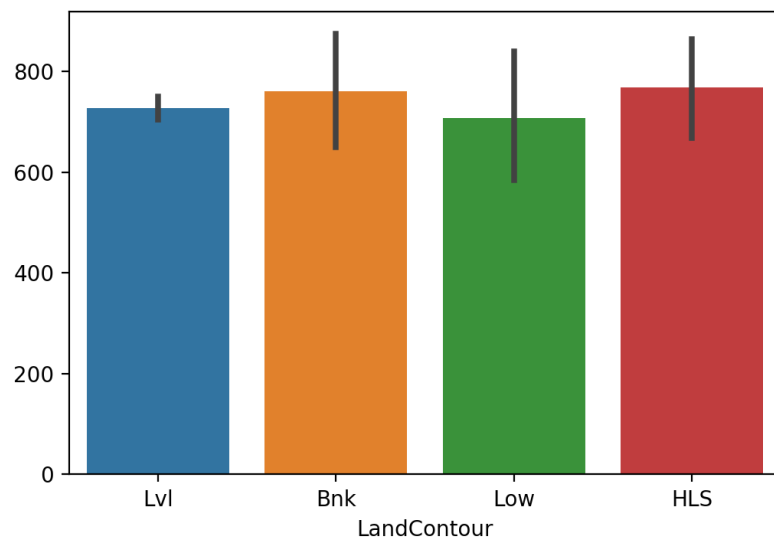

```


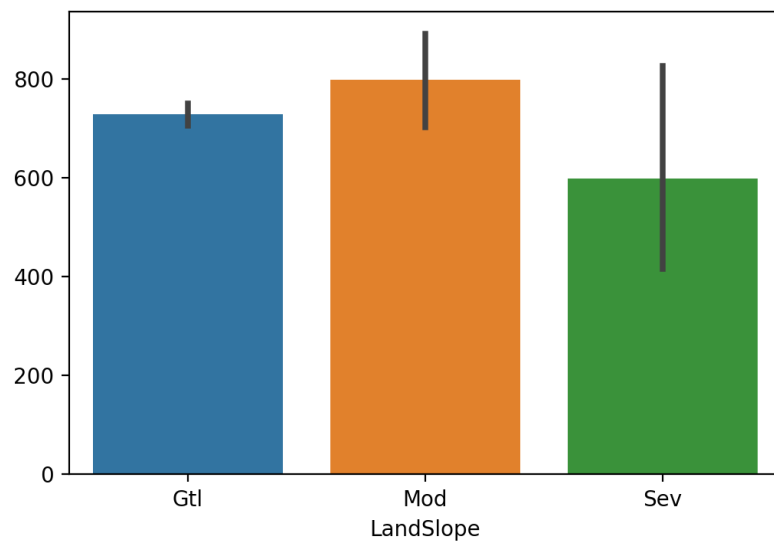
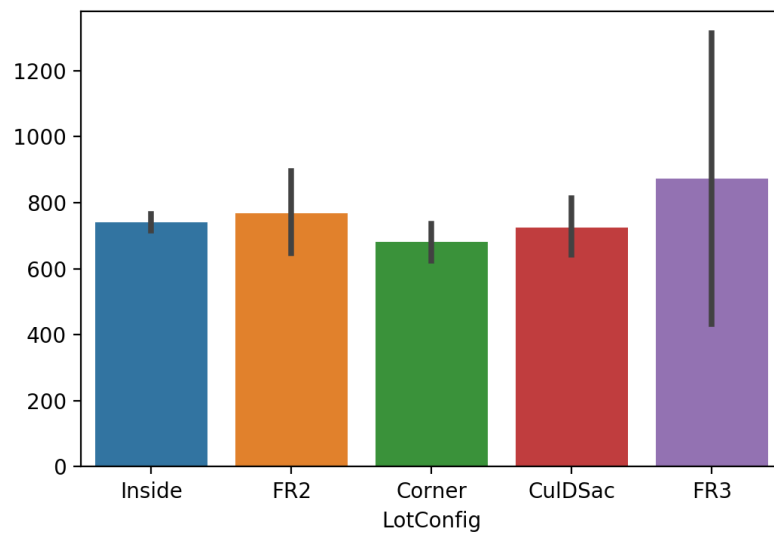
/Users/karinseeder/anaconda3/lib/python3.7/site-packages/matplotlib/pyplot.py:514: RuntimeWarning: More than 20 figures have been opened. Figures created through the pyplot interface (`matplotlib.pyplot.figure`) are retained until explicitly closed and may consume too much memory. (To control this warning, see the rcParam `figure.max_open_warning`).

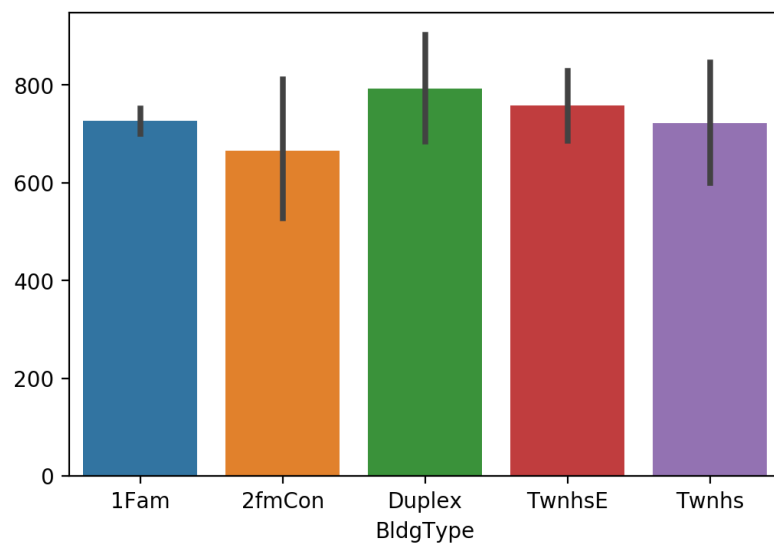
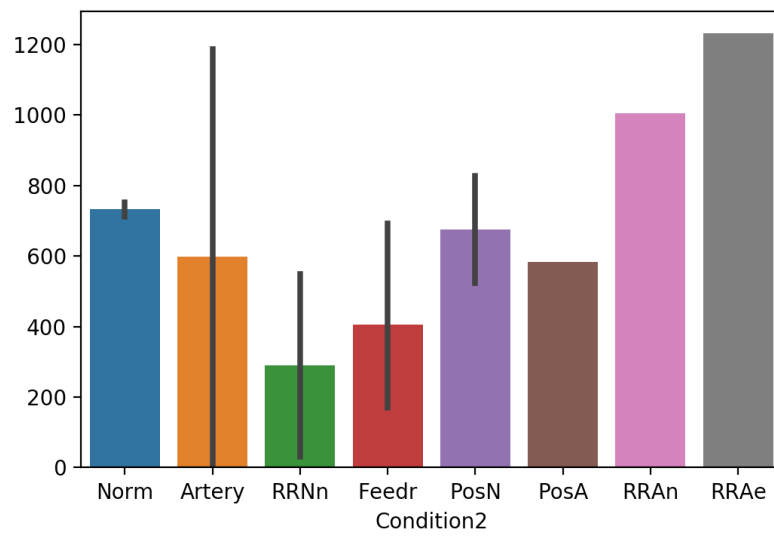

max_open_warning, RuntimeWarning)




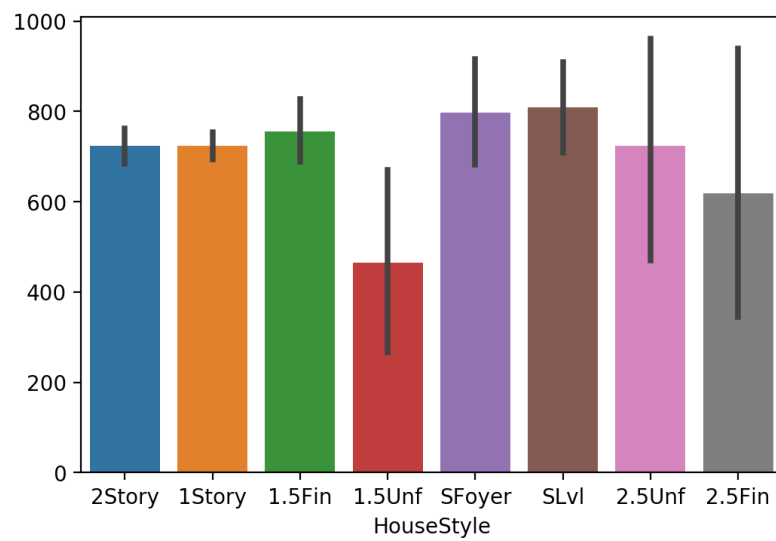


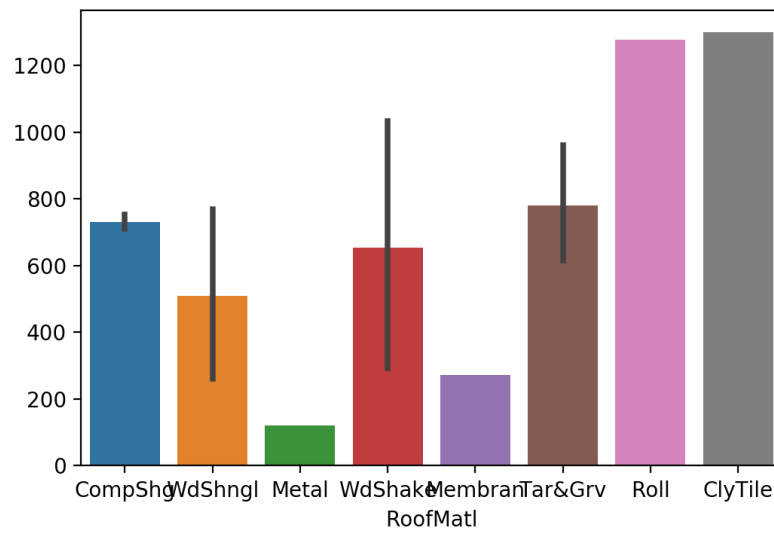
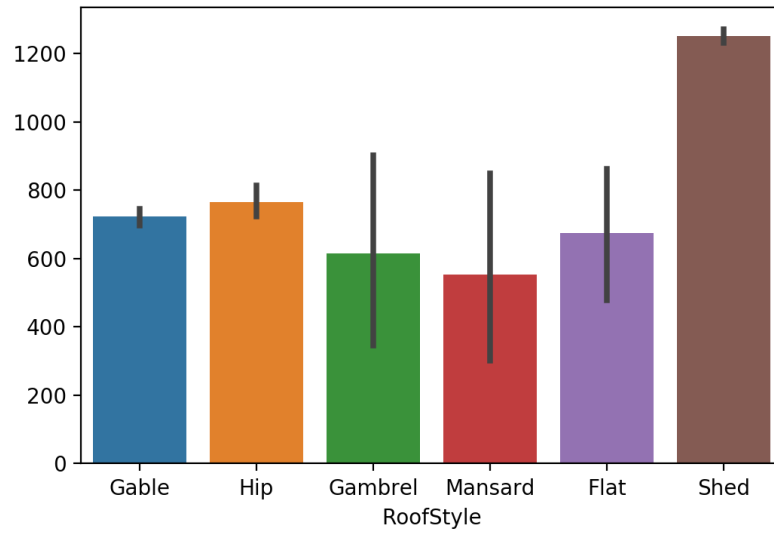



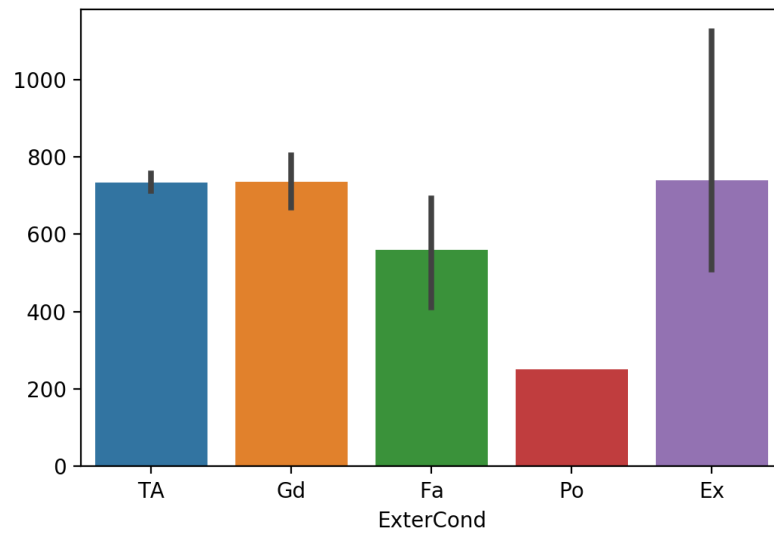
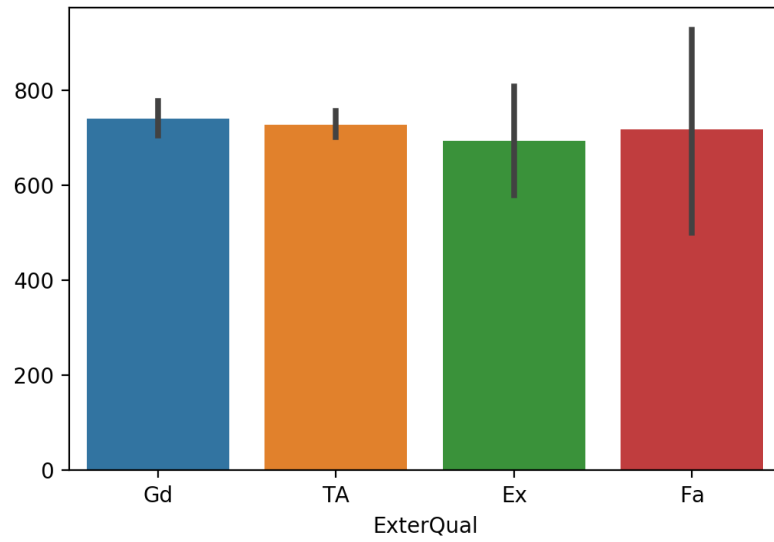



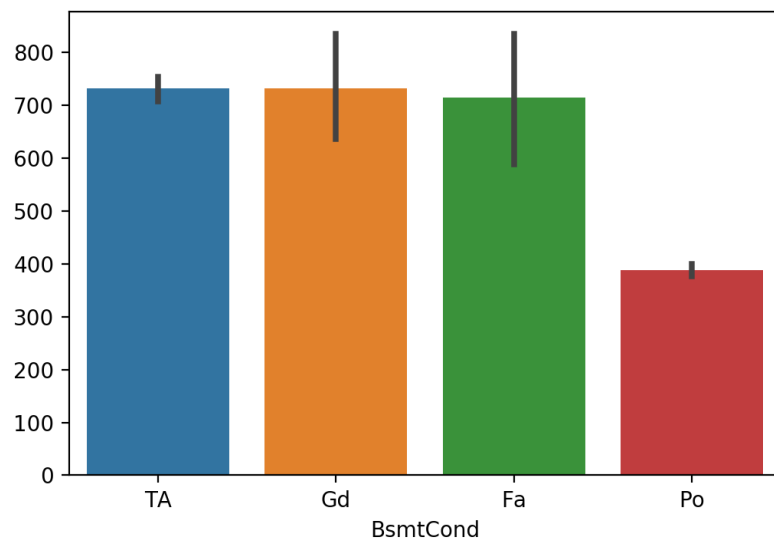

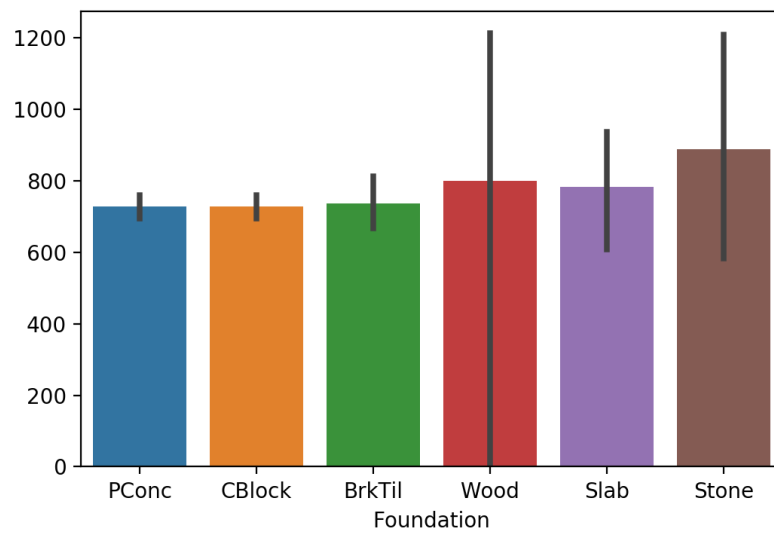




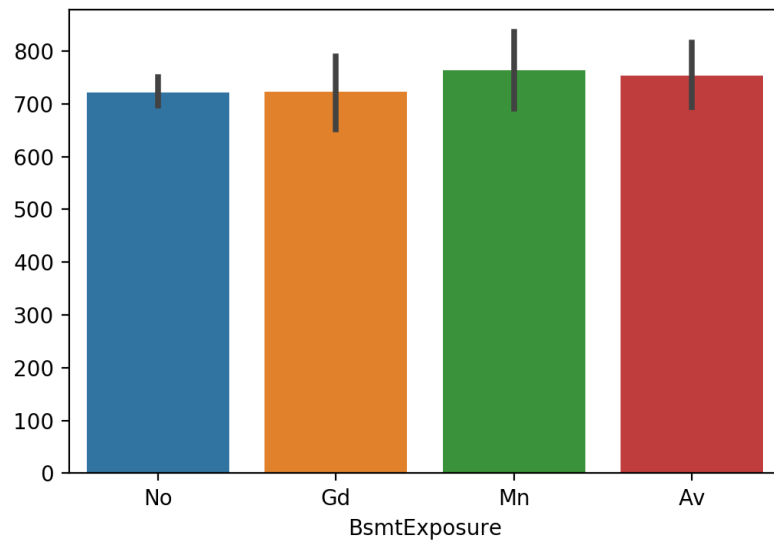























































































```
for i, x in enumerate(df.dtypes):
    if x == 'float64':
        plt.figure(i)
        seaborn.scatterplot(x = df.SalePrice ,y = df[df.columns[i]])
```



2 - Análise de correlações ($r_{xi,xj}$)

2.1 Correlograma (1,5)


```
df_corr = df.corr()

import matplotlib.pyplot as plt
fig, ax = plt.subplots(figsize=(20, 12))
```



```
# mask
mask = np.triu(np.ones_like(df_corr, dtype=np.bool))
# adjust mask and df
mask = mask[1:, :-1]
corr = df_corr.iloc[1:, :-1].copy()
# plot heatmap
seaborn.heatmap(corr, mask=mask, annot=True, fmt=".2f", cmap='Blues',
                 vmin=-1, vmax=1, cbar_kws={'shrink': .8})

# yticks
plt.yticks(rotation=0)
plt.show()
```



2.2 Análise sobre correlações significativas (1,5)

Correlações com a variável resposta

```
df_corr().loc['SalePrice']
```

```
MSSubClass      -0.084284
LotFrontage      0.351799
LotArea          0.263843
OverallQual      0.790982
OverallCond     -0.077856
YearBuilt        0.522897
YearRemodAdd     0.507101
MasVnrArea       0.477493
BsmtFinSF1       0.386420
BsmtFinSF2     -0.011378
BsmtUnfSF        0.214479
TotalBsmtSF      0.613581
1stFlrSF         0.605852
2ndFlrSF         0.319334
LowQualFinSF    -0.025606
GrLivArea        0.708624
BsmtFullBath     0.227122
BsmtHalfBath    -0.016844
FullBath         0.560664
HalfBath         0.284108
BedroomAbvGr    0.168213
KitchenAbvGr    -0.135907
TotRmsAbvGrd    0.533723
Fireplaces       0.466929
GarageYrBlt     0.486362
GarageCars       0.640409
GarageArea       0.623431
WoodDeckSF      0.324413
OpenPorchSF     0.315856
EnclosedPorch   -0.128578
3SsnPorch        0.044584
ScreenPorch     0.111447
PoolArea         0.092404
MiscVal         -0.021190
MoSold           0.046432
YrSold          -0.028923
SalePrice        1.000000
Name: SalePrice, dtype: float64
```

Podemos notar que as variáveis com maior correlação são:

-OverallQual 0.790982 - Faz todo sentido, dado que aqui é uma nota que dão para o imóvel

-TotalBsmtSF 0.613581 - Talvez faça sentido o tamanho do sótão, pois existe a possibilidade de virar um espaço para alugar

-1stFlrSF 0.605852 - Aqui quanto maior a metragem do primeiro andar, maior a área comum, logo um valor maior

-GrLivArea 0.708624 - Idem ao de cima, tamanho maior da área externa de convivência

-GarageCars 0.640409 - Quantidade de carro que cabem na garagem valorizam o imóvel

-GarageArea 0.623431 - Muita correlação com o de cima (0.88)

3 - Desenvolvimento de modelo de Regressão utilizando Regressão Linear com o método de mínimos quadrados ordinários.

Apresente as características do desenvolvimento: amostras, medidas de avaliação do modelo...

Primeiramente realizamos um tratamento nas variáveis categóricas e nos missings

```
categorical_data = ['MSSubClass', 'MSZoning', 'Street', 'Alley', 'LotShape', 'LandContour', 'Utilities', 'LotConfig',
                    'LandSlope', 'Neighborhood', 'Condition1', 'Condition2', 'BldgType', 'HouseStyle', 'RoofStyle', 'RoofMatl', 'Exterior1st',
                    'Exterior2nd', 'MasVnrType', 'ExterQual', 'ExterCond', 'Foundation', 'BsmtQual', 'BsmtCond', 'BsmtExposure', 'BsmtFinType1',
                    'BsmtFinType2', 'Heating', 'HeatingQC', 'CentralAir', 'Electrical', 'KitchenQual', 'Functional', 'FireplaceQu', 'GarageType',
                    'GarageFinish', 'GarageQual', 'GarageCond', 'PavedDrive', 'PoolQC', 'Fence', 'MiscFeature', 'SaleType', 'SaleCondition']

num_data = ['LotFrontage', 'LotArea', 'OverallQual', 'OverallCond', 'MasVnrArea', 'BsmtFinSF1', 'BsmtFinSF2', 'BsmtUnfSF',
            'TotalBsmtSF', '1stFlrSF', '2ndFlrSF', 'LowQualFinSF', 'GrLivArea', 'BsmtFullBath', 'BsmtHalfBath', 'FullBath', 'HalfBath',
            'BedroomAbvGr', 'KitchenAbvGr', 'TotRmsAbvGrd', 'Fireplaces', 'GarageCars', 'GarageArea', 'WoodDeckSF', 'OpenPorchSF',
            'EnclosedPorch', '3SsnPorch', 'ScreenPorch', 'PoolArea', 'MiscVal']

drop_data = ['id']

date_data = ['YearBuilt', 'YearRemodAdd', 'GarageYrBlt', 'MoSold', 'YrSold']

Y = df.SalePrice

X_cat_df = pd.get_dummies(df[categorical_data].fillna('NA'))

X_num_data = df[num_data].fillna(0)

df['garageTime'] = df.YrSold - df.GarageYrBlt

df['timeToSell'] = df.YrSold - df.YearBuilt

X = pd.concat([X_cat_df, X_num_data, df.garageTime.fillna(0), df.timeToSell.fillna(0)], axis = 1)
```

Importando as bibliotecas de ML

```
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.linear_model import LinearRegression
from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import r2_score, mean_absolute_error, mean_squared_error
```

Aplicamos uma taxa de amostragem para teste de 30%, deixando 70% para treino

```
X_train, X_test, y_train, y_test = train_test_split(X, Y, test_size=0.3, random_state=1234)
```

Para selecionar as variáveis mais relevantes fitamos uma random forest e selecionamos as 20 variáveis mais relevantes.

```
rf = RandomForestRegressor(n_estimators=900, n_jobs=-1, max_depth=5)
rf.fit(X_train, y_train)
best_feat = list(pd.DataFrame(rf.feature_importances_, index=X_train.columns, columns=['Importance']).sort_values('Importance').index)
```

```
['OverallQual',
 'GrLivArea',
 '2ndFlrSF',
 'TotalBsmtSF',
 '1stFlrSF',
 'BsmtFinSF1',
 'LotArea',
 'GarageArea',
 'TotRmsAbvGrd',
 'FullBath',
 'GarageCars',
 'timeToSell',
 'MasVnrArea',
 'WoodDeckSF',
 'LotFrontage',
 'Fireplaces',
```

```
'GarageType_Detchd',
'OpenPorchSF',
'KitchenQual_Gd',
'FireplaceQu_NA']
```

```
lr = LinearRegression()
lr.fit(X_train[best_feat], y_train)
```

```
LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None,
normalize=False)
```

```
print('Test R2 =', r2_score(y_test, lr.predict(X_test[best_feat])))
print('Train R2 =', r2_score(y_train, lr.predict(X_train[best_feat])))
```

```
Test R2 = 0.8226256109446395
Train R2 = 0.7782552445955451
```

```
print('Test MAE =', mean_absolute_error(y_test, lr.predict(X_test[best_feat])))
print('Train MAE =', mean_absolute_error(y_train, lr.predict(X_train[best_feat])))
```

```
Test MAE = 22034.64190126019
Train MAE = 23742.679130796634
```

```
print('Test RMSE =', np.sqrt(mean_squared_error(y_test, lr.predict(X_test[best_feat]))))
print('Train RMSE =', np.sqrt(mean_squared_error(y_train, lr.predict(X_train[best_feat]))))
```

```
Test RMSE = 29774.417592378973
Train RMSE = 38998.84589651374
```

Importando as bibliotecas para o teste de ANOVA

```
from statsmodels.api import OLS, add_constant, graphics
from scipy import stats
```

Avaliando a significancia dos coeficientes e analise de residuos:

```
X2_train = add_constant(X_train[best_feat])
est = OLS(y_train, X2_train)
est2 = est.fit()
print(est2.summary())

seaborn.residplot(lr.predict(X_test[best_feat]), y_test, lowess=True,
                  line_kws={'color': 'red', 'lw': 1, 'alpha': 1})
plt.xlabel("Fitted values")
plt.title('Residual plot')
```

```

                        OLS Regression Results
=====
Dep. Variable:          SalePrice   R-squared:
0.778
Model:                  OLS        Adj. R-squared:
0.774
Method:                 Least Squares   F-statistic:
175.7
Date:                   Thu, 29 Jul 2021   Prob (F-statistic):
2.49e-310
Time:                   00:29:21         Log-Likelihood:
-12254.
No. Observations:      1022   AIC:
2.455e+04
Df Residuals:          1001   BIC:
2.465e+04
Df Model:              20
Covariance Type:       nonrobust
=====
[0.025    0.975]      coef      std err      t      P>|t|
-----
const              -7.423e+04   1.16e+04   -6.421    0.000
-9.69e+04  -5.15e+04
OverallQual        2.219e+04   1503.530   14.759    0.000
1.92e+04   2.51e+04
GrLivArea          41.9901    26.280    1.598    0.110
-9.579    93.559
2ndFlrSF           -6.0037    26.479   -0.227    0.821
-57.964    45.957
TotalBsmtSF        8.3485     5.604    1.490    0.137
```

-2.648	19.345				
1stFlrSF		-2.3063	26.917	-0.086	0.932
-55.127	50.515				
BsmtFinSF1		14.9089	3.397	4.389	0.000
8.242	21.575				
LotArea		0.3325	0.146	2.284	0.023
0.047	0.618				
GarageArea		10.8056	12.868	0.840	0.401
-14.446	36.058				
TotRmsAbvGrd		2143.3835	1407.089	1.523	0.128
-617.799	4904.566				
FullBath		-400.3125	3426.292	-0.117	0.907
-7123.852	6323.227				
GarageCars		1.142e+04	3890.672	2.934	0.003
3781.777	1.91e+04				
timeToSell		-222.2222	66.547	-3.339	0.001
-352.810	-91.635				
MasVnrArea		32.0146	7.761	4.125	0.000
16.785	47.244				
WoodDeckSF		40.4396	10.432	3.877	0.000
19.969	60.911				
LotFrontage		9.2440	37.745	0.245	0.807
-64.824	83.312				
Fireplaces		7901.7778	4892.097	1.615	0.107
-1698.164	1.75e+04				
GarageType_Detchd		-5080.0312	3496.892	-1.453	0.147
-1.19e+04	1782.048				
OpenPorchSF		6.2005	19.851	0.312	0.755
-32.753	45.154				
KitchenQual_Gd		-2579.8473	2997.803	-0.861	0.390
-8462.546	3302.851				
FireplaceQu_NA		3203.0148	6216.780	0.515	0.607
-8996.401	1.54e+04				

Omnibus:	445.486	Durbin-Watson:
2.117		
Prob(Omnibus):	0.000	Jarque-Bera (JB):
57093.269		
Skew:	-0.952	Prob(JB):
0.00		
Kurtosis:	39.567	Cond. No.
1.47e+05		


Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 1.47e+05. This might indicate that there are strong multicollinearity or other numerical problems.

/Users/karinseeder/anaconda3/lib/python3.7/site-packages/statsmodels/tsa/tsatools.py:142: FutureWarning: In a future version of pandas all arguments of concat except for the argument 'objs' will be keyword-only
x = pd.concat(x[:, :order], 1)

Text(0.5, 1.0, 'Residual plot')




```
residuals = y_test - lr.predict(X_test[best_feat])
residuals

plt.figure(figsize=(7,7))
```

```
stats.probplot(residuals, dist="norm", plot=plt)
plt.title("Normal Q-Q Plot")
```


```
Text(0.5, 1.0, 'Normal Q-Q Plot')
```



```
model_norm_residuals_abs_sqrt=np.sqrt(np.abs(residuals))

plt.figure(figsize=(7,7))
seaborn.regplot(lr.predict(X_test[best_feat]), model_norm_residuals_abs_sqrt,
                 scatter=True,
                 lowess=True,
                 line_kws={'color': 'red', 'lw': 1, 'alpha': 0.8})
plt.ylabel("Standarized residuals")
plt.xlabel("Fitted value")
plt.title('Scale-Location plot')
```

```
Text(0.5, 1.0, 'Scale-Location plot')
```



```
from scipy import stats
shapiro_test = stats.shapiro(residuals)
print(shapiro_test)
```

```
ShapiroResult(statistic=0.9504468441009521,
pvalue=6.039371086563605e-11)
```

Conclusões

Analizando o R^2 podemos verificar que o modelo explica 78% da variável resposta, ou seja, as variáveis escolhidas explicam 78% do preço da casa.

Avaliando a Prob-F, dado que ela é menor que 0.05 podemos rejeitar a hipótese H_0 de que a regressão linear não existe.

Na métrica de Durbin-Watson também podemos concluir que não há correlação entre os resíduos.

Apenas com estas métricas poderíamos concluir que temos um bom modelo, porém ao continuar nossas análises podemos notar que tanto MAE e RMSE apresentam valores significantes e a diferença entre o RMSE e MAE nos permite supor que possa existir algum outlier e que os erros não são uniformes.

Pela análise de resíduo podemos verificar graficamente que os erros não são uniformemente distribuídos e que apresentam um leve formato de U. Aplicando o teste de Shapiro com o resultado de $p\text{-value} < 0.05$, isso nos leva a aceitar a hipótese nula de que os erros não seguem uma distribuição normal. O mesmo resultado pode ser obtido pela Prob Omnibus apresentada na tabela OLS Regression Results.

Portanto, podemos concluir que a regressão linear não é a melhor escolha.