

E2 - Classificação com Regressão Logística (Análise crédito)

"South German Credit" é uma base com informações de 1000 créditos (700 good and 300 bad) com 20 variáveis preditoras.

Objetivo: desenvolver um modelo de predição do risco de crédito (0 = bom; 1 = ruim)

- Importando as bibliotecas

```
import pandas as pd
import numpy as np
import statsmodels
import seaborn
from matplotlib import pyplot as plt
pd.options.display.max_columns = 100
pd.options.mode.chained_assignment = None # default='warn'
```

```
/Users/karinseeder/anaconda3/lib/python3.7/site-
packages/pandas/compat/_optional.py:138: UserWarning: Pandas requires
version '2.7.0' or newer of 'numexpr' (version
'2.6.8' currently installed).
  warnings.warn(msg, UserWarning)
```

- Importando a base de dados

```
df = pd.read_csv('base_2sgc.csv')
print(df.shape)
```

```
(1000, 21)
```

- Visualizando uma amostra da base

```
df.head()
```

	status	duration	credit_history	purpose	amount	savings	employment_duration	installment_rate	credit_risk
0	no checking account	18	all credits at this bank paid back duly	car (used)	1049	unknown/no savings account	< 1 yr	< 20	for credit
1	no checking account	9	all credits at this bank paid back duly	others	2799	unknown/no savings account	1 <= ... < 4 yrs	25 <= ... < 35	refused
2	... < 0 DM	12	no credits taken/all credits paid back duly	retraining	841	... < 100 DM	4 <= ... < 7 yrs	25 <= ... < 35	for credit
3	no checking account	12	all credits at this bank paid back duly	others	2122	unknown/no savings account	1 <= ... < 4 yrs	20 <= ... < 25	refused
4	no checking account	12	all credits at this bank paid back duly	others	2171	unknown/no savings account	1 <= ... < 4 yrs	< 20	refused

1 - Análise descritiva de variáveis

1.1 Estatísticas descritivas: frequência, tabelas cruzadas, média (\bar{x}), desvio padrão (s), quartis ($Q1$, \bar{x} , $Q3$) (2,0)

- Avaliando o tipo das variáveis na base

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1000 entries, 0 to 999
Data columns (total 21 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   status                                1000 non-null   object
1   duration                              1000 non-null   int64
2   credit_history                        1000 non-null   object
3   purpose                              1000 non-null   object
4   amount                               1000 non-null   int64
5   savings                              1000 non-null   object
6   employment_duration                  1000 non-null   object
7   installment_rate                     1000 non-null   object
8   personal_status_sex                  1000 non-null   object
9   other_debtors                        1000 non-null   object
10  present_residence                    1000 non-null   object
11  property                             1000 non-null   object
12  age                                   1000 non-null   int64
13  other_installment_plans               1000 non-null   object
14  housing                              1000 non-null   object
15  number_credits                       1000 non-null   object
16  job                                   1000 non-null   object
17  people_liable                        1000 non-null   object
18  telephone                           1000 non-null   object
19  foreign_worker                       1000 non-null   object
20  credit_risk                          1000 non-null   object
dtypes: int64(3), object(18)
memory usage: 164.2+ KB
```

- Análise descritiva para variáveis numéricas

```
df.describe()
```

	duration	amount	age
count	1000.000000	1000.000000	1000.000000
mean	20.903000	3271.24800	35.54200
std	12.058814	2822.75176	11.35267
min	4.000000	250.00000	19.00000
25%	12.000000	1365.50000	27.00000
50%	18.000000	2319.50000	33.00000
75%	24.000000	3972.25000	42.00000
max	72.000000	18424.00000	75.00000

- Análise descritiva das variáveis categóricas

```
df.describe(include=object)
```

	status	credit_history	purpose	savings	employment_duration	installment_rate	personal_status
count	1000	1000	1000	1000	1000	1000	1000
unique	4	5	10	5	5	4	4
top	... >= 200 DM / salary for at least 1 year	no credits taken/all credits paid back duly	furniture/equipment	unknown/no savings account	1 <= ... < 4 yrs	< 20	male : married/wic

	status	credit_history	purpose	savings	employment_duration	installment_rate	personal_status
freq	394	530	280	603	339	476	548

```
plt.rcParams["figure.figsize"] = (20,10)
for i, x in enumerate(df.dtypes):
    if x == 'object':
        #print(pd.crosstab(index=df[df.columns[i]], columns='freq', dropna=False))
        print(pd.crosstab(df[df.columns[i]], df['credit_risk']).apply(lambda r: r/r.sum(), axis=0))
        print('')
```

```
credit_risk
status
... < 0 DM          0.350000  0.234286
... >= 200 DM / salary for at least 1 year 0.153333  0.497143
0<= ... < 200 DM    0.046667  0.070000
no checking account 0.450000  0.198571
```

```
credit_risk
credit_history
all credits at this bank paid back duly 0.166667  0.347143
critical account/other credits elsewhere 0.093333  0.030000
delay in paying off in the past         0.083333  0.021429
existing credits paid back duly till now 0.093333  0.085714
no credits taken/all credits paid back duly 0.563333  0.515714
```

```
credit_risk
purpose
business          0.016667  0.010000
car (new)          0.056667  0.122857
car (used)         0.193333  0.175714
domestic appliances 0.026667  0.020000
furniture/equipment 0.206667  0.311429
others            0.296667  0.207143
radio/television   0.013333  0.011429
repairs           0.073333  0.040000
retraining         0.113333  0.090000
vacation          0.003333  0.011429
```

```
credit_risk
savings
... < 100 DM        0.113333  0.098571
... >= 1000 DM      0.106667  0.215714
100 <= ... < 500 DM 0.036667  0.074286
500 <= ... < 1000 DM 0.020000  0.060000
unknown/no savings account 0.723333  0.551429
```

```
credit_risk
employment_duration
1 <= ... < 4 yrs    0.346667  0.335714
4 <= ... < 7 yrs    0.130000  0.192857
< 1 yr             0.233333  0.145714
>= 7 yrs           0.213333  0.270000
unemployed         0.076667  0.055714
```

```
credit_risk
installment_rate
20 <= ... < 25      0.150000  0.160000
25 <= ... < 35      0.206667  0.241429
< 20               0.530000  0.452857
>= 35              0.113333  0.145714
```

```
credit_risk
personal_status_sex
female : non-single or male : single 0.363333  0.287143
female : single                      0.083333  0.095714
male : divorced/separated            0.066667  0.042857
male : married/widowed               0.486667  0.574286
```

```
credit_risk
other_debtors
co-applicant 0.060000  0.032857
guarantor    0.033333  0.060000
none         0.906667  0.907143
```

```
credit_risk
present_residence
1 <= ... < 4 yrs    0.323333  0.301429
4 <= ... < 7 yrs    0.143333  0.151429
< 1 yr            0.120000  0.134286
>= 7 yrs          0.413333  0.412857
```

```
credit_risk
bad      good
```

property			
building soc. savings agr./life insurance	0.340000	0.328571	
car or other	0.236667	0.230000	
real estate	0.223333	0.124286	
unknown / no property	0.200000	0.317143	

credit_risk	bad	good
other_installment_plans		
bank	0.190000	0.117143
none	0.746667	0.842857
stores	0.063333	0.040000

credit_risk	bad	good
housing		
for free	0.233333	0.155714
own	0.146667	0.090000
rent	0.620000	0.754286

credit_risk	bad	good
number_credits		
1	0.666667	0.618571
2-3	0.306667	0.344286
4-5	0.020000	0.031429
>= 6	0.006667	0.005714

credit_risk		bad	good
job			
manager/self-empl./highly qualif. employee	0.170000	0.138571	
skilled employee/official	0.620000	0.634286	
unemployed/unskilled - non-resident	0.023333	0.021429	
unskilled - resident	0.186667	0.205714	

credit_risk	bad	good
people_liable		
0 to 2	0.846667	0.844286
3 or more	0.153333	0.155714

credit_risk		bad	good
telephone			
no		0.623333	0.584286
yes (under customer name)		0.376667	0.415714

credit_risk	bad	good
foreign_worker		
no	0.986667	0.952857
yes	0.013333	0.047143

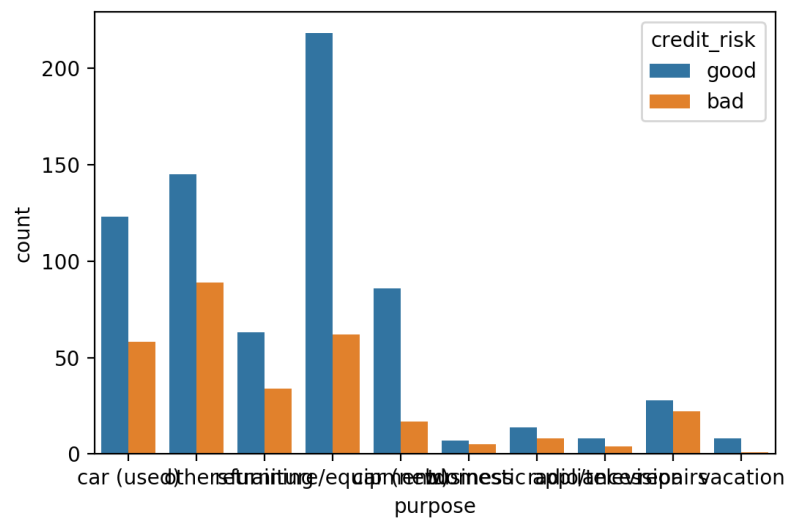
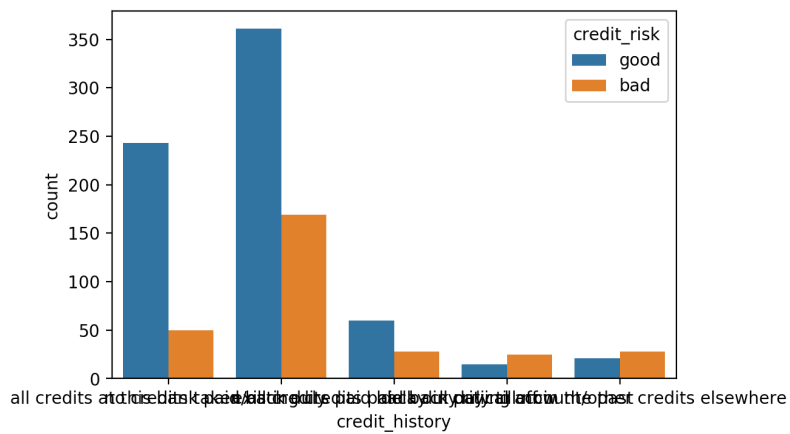

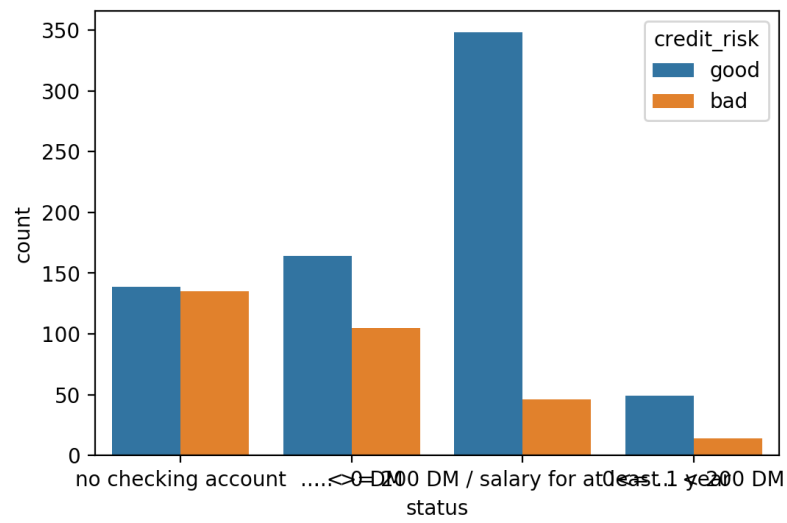
credit_risk	bad	good
credit_risk		
bad	1.0	0.0
good	0.0	1.0

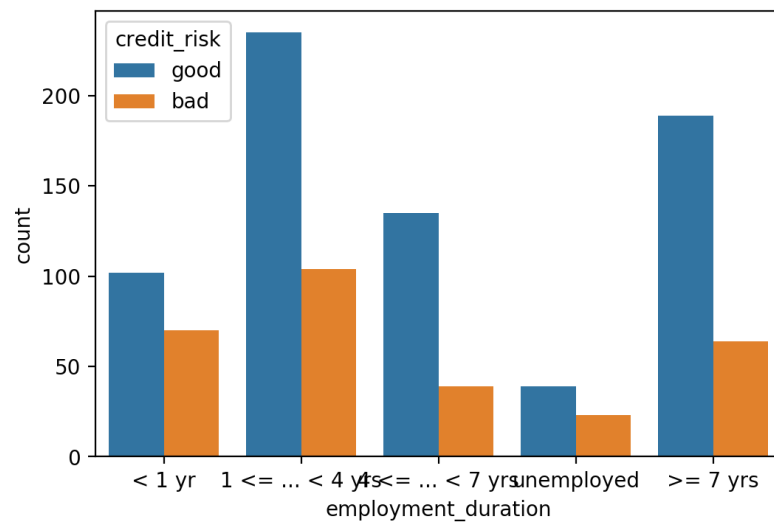
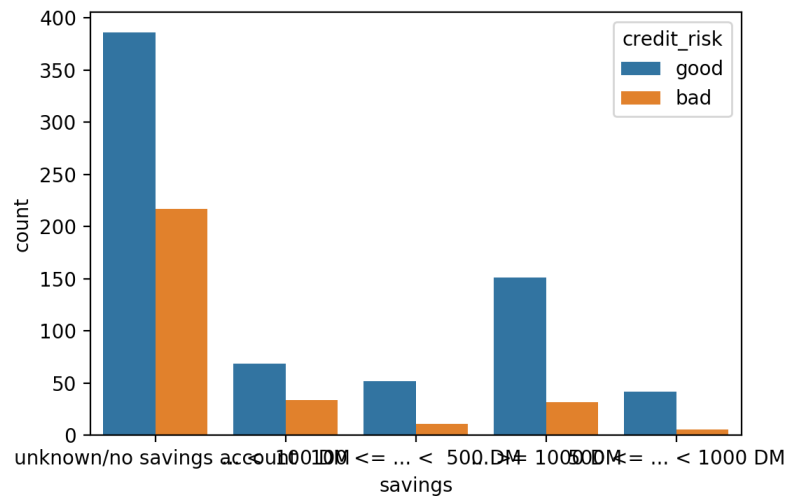

1.2 Gráficos como: Gráficos de columnas, BoxPlot por categorías (2,0)

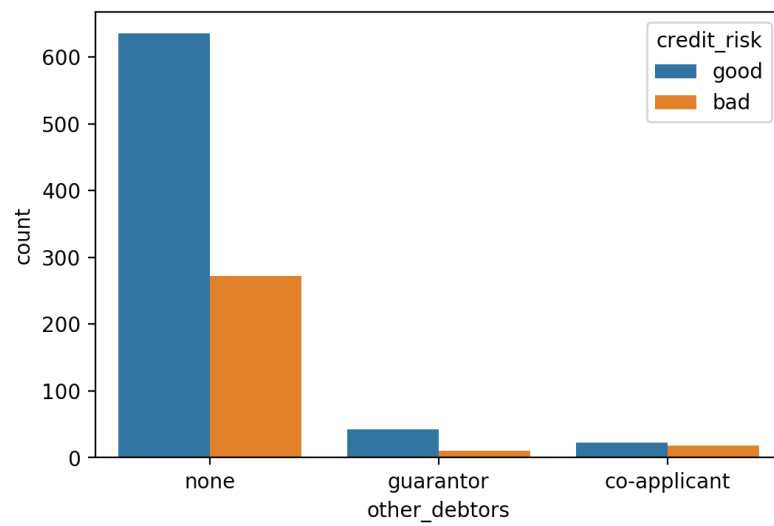
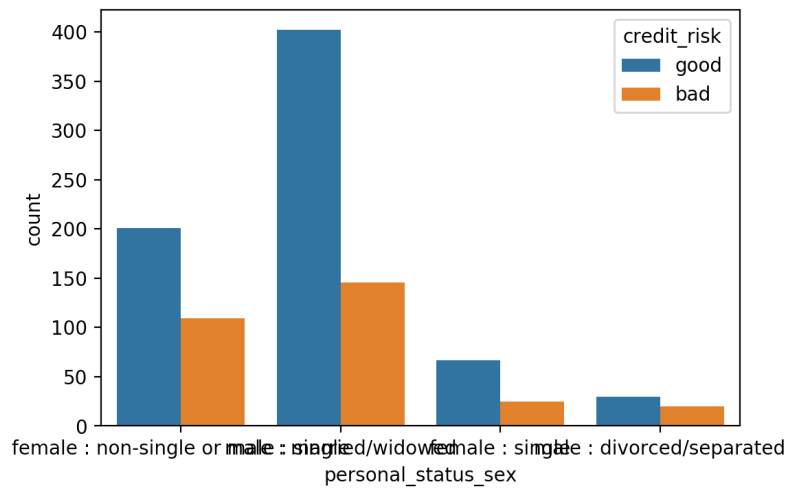

```
import seaborn as sns


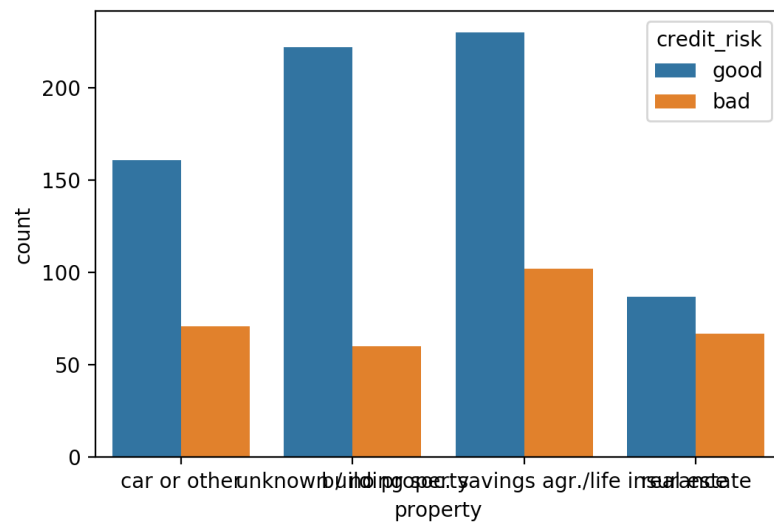

for i, x in enumerate(df.dtypes):
    if x == 'int64' or x == 'float64':
        plt.figure(i)
        sns.boxplot(data = df, y = df[df.columns[i]], hue = df['credit_risk'])
    elif x == 'object':
        plt.figure(i)
        sns.countplot(x=df[df.columns[i]], hue = 'credit_risk', data=df)
```


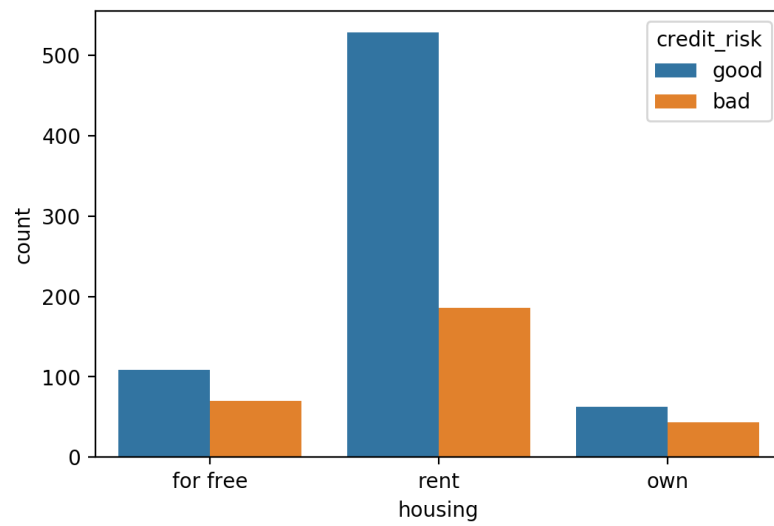
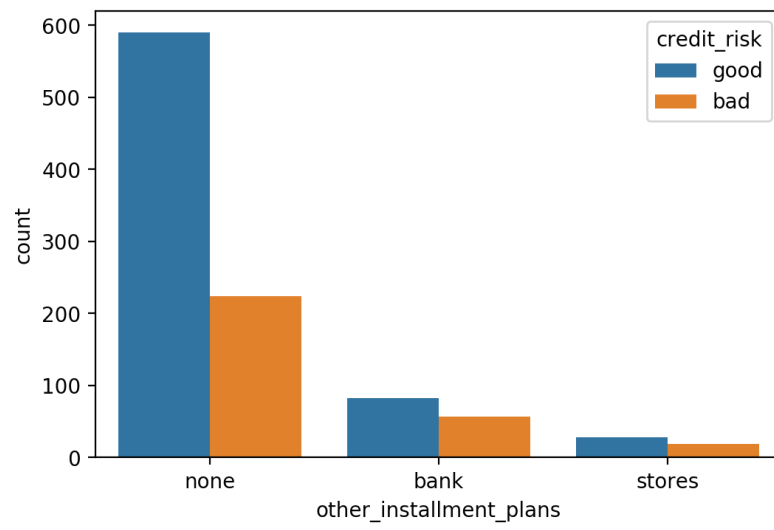
```
/Users/karinseeder/anaconda3/lib/python3.7/site-
packages/matplotlib/pyplot.py:514: RuntimeWarning: More than 20
figures have been opened. Figures created through the pyplot interface
(`matplotlib.pyplot.figure`) are retained until explicitly closed and
may consume too much memory. (To control this warning, see the rcParam
`figure.max_open_warning`).
  max_open_warning, RuntimeWarning)
```

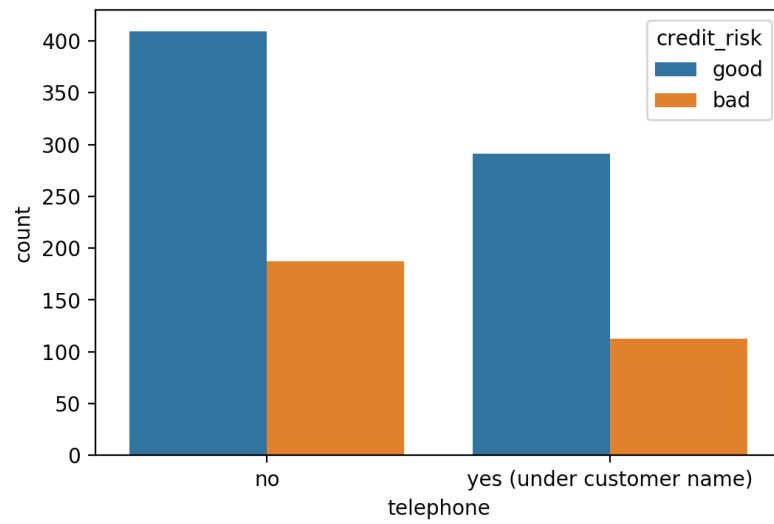
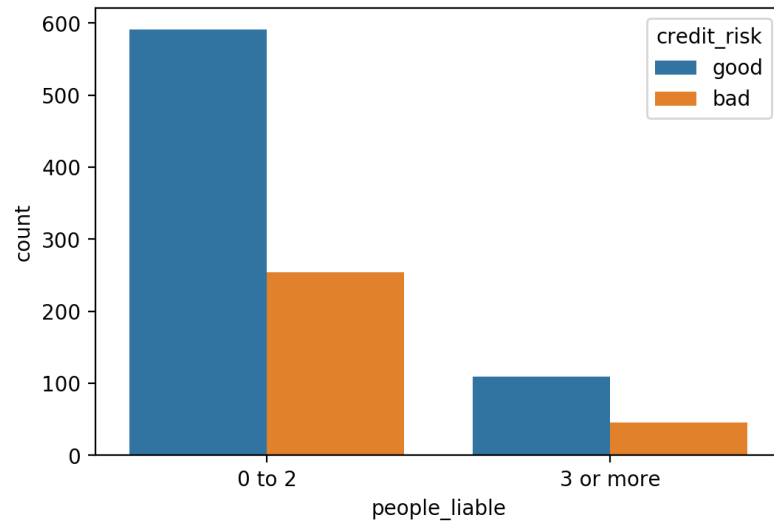
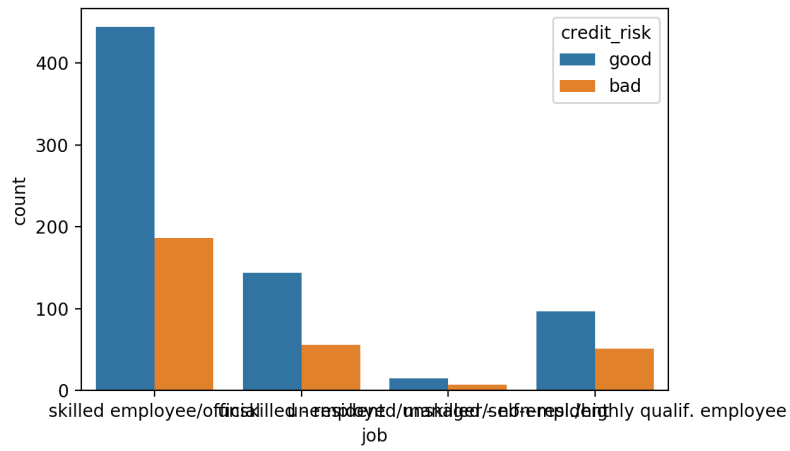



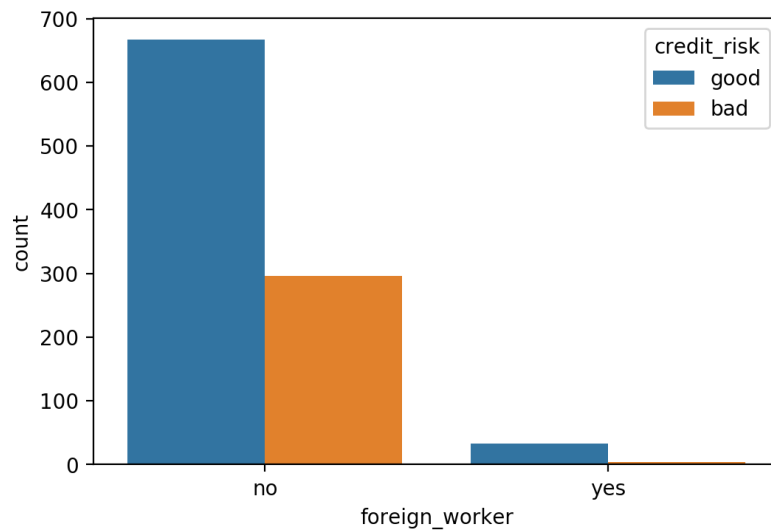












2 - Desenvolvimento de modelo de Classificação utilizando Regressão Logística (binomial).

- Primeiramente realizamos um tratamento nas variáveis

```
cols = df.columns
num_data = list(df._get_numeric_data().columns)
categorical_data = list(set(cols) - set(num_data) - set(['credit_risk']))
Y = df.credit_risk
X_cat = df[categorical_data]
X_num_data = df[num_data]

factorize_cols = ['status', 'savings', 'employment_duration', 'installment_rate', 'present_residence', 'number_credits', 'people_liable']
dummies_cols = list(set(categorical_data) - set(factorize_cols))
```

```
cat_status = ["no checking account", "... < 0 DM", "0<= ... < 200 DM", "... >= 200 DM / salary for at least 1 year"]
cat_savings = ["unknown/no savings account", "... < 100 DM", "100 <= ... < 500 DM", "500 <= ... < 1000 DM", "... >= 1000 DM"]
cat_employment_duration = ["unemployed", "< 1 yr", "1 <= ... < 4 yrs", "4 <= ... < 7 yrs", ">= 7 yrs"]
cat_installment_rate = [">= 35", "25 <= ... < 35", "20 <= ... < 25", "< 20"]
cat_present_residence = ["< 1 yr", "1 <= ... < 4 yrs", "4 <= ... < 7 yrs", ">= 7 yrs"]
cat_number_credits = ["1", "2-3", "4-5", ">= 6"]
cat_people_liable = ["3 or more", "0 to 2"]

X_cat['status'] = pd.factorize(pd.Categorical(X_cat['status'], categories = cat_status))[0]
X_cat['savings'] = pd.factorize(pd.Categorical(X_cat['savings'], categories = cat_savings))[0]
X_cat['employment_duration'] = pd.factorize(pd.Categorical(X_cat['employment_duration'], categories = cat_employment_duration))[0]
X_cat['installment_rate'] = pd.factorize(pd.Categorical(X_cat['installment_rate'], categories = cat_installment_rate))[0]
X_cat['present_residence'] = pd.factorize(pd.Categorical(X_cat['present_residence'], categories = cat_present_residence))[0]
X_cat['number_credits'] = pd.factorize(pd.Categorical(X_cat['number_credits'], categories = cat_number_credits))[0]
X_cat['people_liable'] = pd.factorize(pd.Categorical(X_cat['people_liable'], categories = cat_people_liable))[0]

X_cat = pd.get_dummies(X_cat)

Y = pd.factorize(Y)[0]
```

```
def robust_scaling(df):
    df_robust = df.copy()
    for column in df_robust.columns:
        df_robust[column] = (df_robust[column] - df_robust[column].median()) / (df_robust[column].quantile(0.75) - df_robust[column].quantile(0.25))
    return df_robust

X_num_data = robust_scaling(X_num_data)

X = pd.concat([X_cat, X_num_data], axis = 1)
```

Importando as bibliotecas de ML

```
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import r2_score, mean_absolute_error, mean_squared_error
```

2.1 Descrição da amostra de treino e de teste (0,5)

Aplicamos uma taxa de amostragem para teste de 20%, deixando 80% para treino e estratificado para manter a mesma proporção de bons e mals nas duas amostras.

```
X_train, X_test, y_train, y_test = train_test_split(X, Y, test_size=0.2, random_state=1234, stratify=Y)

log_reg = LogisticRegression()
log = log_reg.fit(X_train, y_train)
y_pred = log.predict(X_train)
y_predtest = log.predict(X_test)
print("train score: ", log.score(X_train, y_train))
print("test score: ", log.score(X_test, y_test))
```

```
train score: 0.7825
test score: 0.775
/Users/karinseeder/anaconda3/lib/python3.7/site-packages/sklearn/linear_model/logistic.py:433: FutureWarning: Default solver will be changed to 'lbfgs' in 0.22. Specify a solver to silence this warning.
  FutureWarning)
```


2.2 Descrição do modelo, coeficientes, significância obtida (2,0)

```
odds = np.exp(log.coef_[0]) / (1 + np.exp(log.coef_[0]))
coeficientes = pd.DataFrame(odds,
                             X_train.columns,
                             columns=['coef'])\
    .sort_values(by='coef', ascending=False)
print(coeficientes)
print(" ")
plt.rcParams["figure.figsize"] = (20,10)
plt.plot(log.coef_.T, 'o', label = "teste")
plt.xticks(range(X_train.shape[1]), X_train.columns, rotation=90)
plt.hlines(0,0, X_train.shape[1])
```

	coef
credit_history_delay in paying off in the past	0.678705
purpose_others	0.669687
purpose_repairs	0.656786
credit_history_critical account/other credits e...	0.648019
personal_status_sex_male : divorced/separated	0.625732
foreign_worker_no	0.605555
purpose_domestic appliances	0.594732
other_debtors_co-applicant	0.583300
duration	0.578365
property_real estate	0.577447
amount	0.561871
telephone_no	0.555429
other_installment_plans_stores	0.539598
housing_for free	0.539195
job_unemployed/unskilled - non-resident	0.537954
housing_own	0.536939
other_debtors_none	0.536445
people_liable	0.529544
personal_status_sex_female : single	0.525398
other_installment_plans_bank	0.518406
personal_status_sex_female : non-single or male...	0.509630
property_car or other	0.509334
number_credits	0.509061
purpose_retraining	0.508028
job_skilled employee/official	0.506645
job_unskilled - resident	0.504119

purpose_car (used)	0.503631
purpose_radio/television	0.491185
property_building soc. savings agr./life insurance	0.489833
employment_duration	0.488828
job_manager/self-empl./highly qualif. employee	0.484633
present_residence	0.484393
purpose_furniture/equipment	0.479797
telephone_yes (under customer name)	0.477778
other_installment_plans_none	0.475346
age	0.474306
credit_history_no credits taken/all credits paid back	0.466320
housing_rent	0.457244
property_unknown / no property	0.456291
installment_rate	0.447800
credit_history_existing credits paid back duly ...	0.441138
purpose_vacation	0.429791
foreign_worker_yes	0.426785
savings	0.422683
other_debtors_guarantor	0.413701
purpose_business	0.391037
personal_status_sex_male : married/widowed	0.372742
status	0.332473
purpose_car (new)	0.307781
credit_history_all credits at this bank paid ba...	0.298801

<matplotlib.collections.LineCollection at 0x7fa31b6e7ef0>




A partir do resultado acima podemos notar que as variáveis mais importantes fazem sentido dado que o histórico de atraso no passado, propósito do empréstimo sem definição podendo ser um emprestimo sem garantia, para fins de reparo tambem podem ser creditos mais arriscado dado que não necessariamente a um valor agregado para aquele montante emprestado. Por fim, crédito em outros bancos já mostra um indviduamento maior do indivíduo, diminuindo sua capacidade de pagamento por já possuir compromisso com outras instituições.

2.3 Curva ROC obtida no treino - análise (1.0)

```
from sklearn.metrics import roc_curve
plt.rcParams["figure.figsize"] = (20,10)
fpr, tpr, thresholds = roc_curve(y_train, y_pred)

def plot_roc(fpr, tpr, thresholds):
    plt.plot(fpr, tpr, label=None)
    plt.plot([0,1], [0,1], 'k--')
    plt.axis([0,1,0,1])
    plt.xlabel("Taxa de Falso positivo")
    plt.ylabel("Taxa de Verdadeiros positivos")

plot_roc(fpr, tpr, thresholds)
plt.show()
```



2.4 Medidas de desempenho para amostra de treino e para a amostra de teste: Matriz confusão, precisão, revocação, $F1$ (2,5)

```
from sklearn.metrics import confusion_matrix, precision_score, recall_score, f1_score
print('treino')
print('matriz de confusão: ')
print(confusion_matrix(y_train, y_pred))
print('Precision: ', precision_score(y_train, y_pred))
print('Recall: ', recall_score(y_train, y_pred))
print('F1: ', f1_score(y_train, y_pred))

print('teste')
print('matriz de confusão: ')
print(confusion_matrix(y_test, y_predtest))
print('Precision: ', precision_score(y_test, y_predtest))
print('Recall: ', recall_score(y_test, y_predtest))
print('F1: ', f1_score(y_test, y_predtest))
```

```
treino
matriz de confusão:
[[502  58]
 [116 124]]
Precision:  0.6813186813186813
Recall:  0.5166666666666667
F1:  0.5876777251184834
teste
matriz de confusão:
[[121  19]
 [ 26  34]]
Precision:  0.6415094339622641
Recall:  0.5666666666666667
F1:  0.6017699115044247
```

Conclusão

Pelos resultados obtidos podemos concluir que o modelo não possui overfitting. Apesar do ajuste apresentar uma acurácia de 78%, a partir dos resultados da matriz de confusão é possível dizer que o modelo não classifica bem os mal pagadores.